






All three panels use prompts evaluable under all four conditions: primary $n = 37$, fresh A $n = 20$, fresh B $n = 18$. Outcome coding: trajectories the harness could not score are excluded. A: simulated attack success rates with 95% Wilson intervals; counts are printed above each primary bar. B: primary-sample mean probe-margin changes versus baseline, with 95% paired percentile-bootstrap intervals. C: mean L14 probe margin versus simulated attack success rate; vertical error bars show 95% Wilson intervals. The 0/96 annotation separately summarizes all available residual readouts. L14 is the primary readout; L18 is diagnostic only, since its probe falls to chance accuracy on the spoofed-declaration distribution ($n=16$, log loss 4.51) even though its downstream response still speaks to role re-entry. $\cos(w_{11}, v_{CAM}) = 0.591$ and $\cos(w_{11}, w_{14}) = 0.76$, so the residual's small L14 shift may partly reflect alignment between the L11 and L14 probe directions; the similar observed attack rates do not follow from that alignment. Similar rates are not an equivalence test.